


B-Trees

COP 3503
Fall 2021


Department of Computer Science
University of Central Florida
Dr. Steinberg




Introduction

- Binary Search Trees – Trees where a node has only at most two children
- Examples:

0
1
2
3



0
1
2
3
4

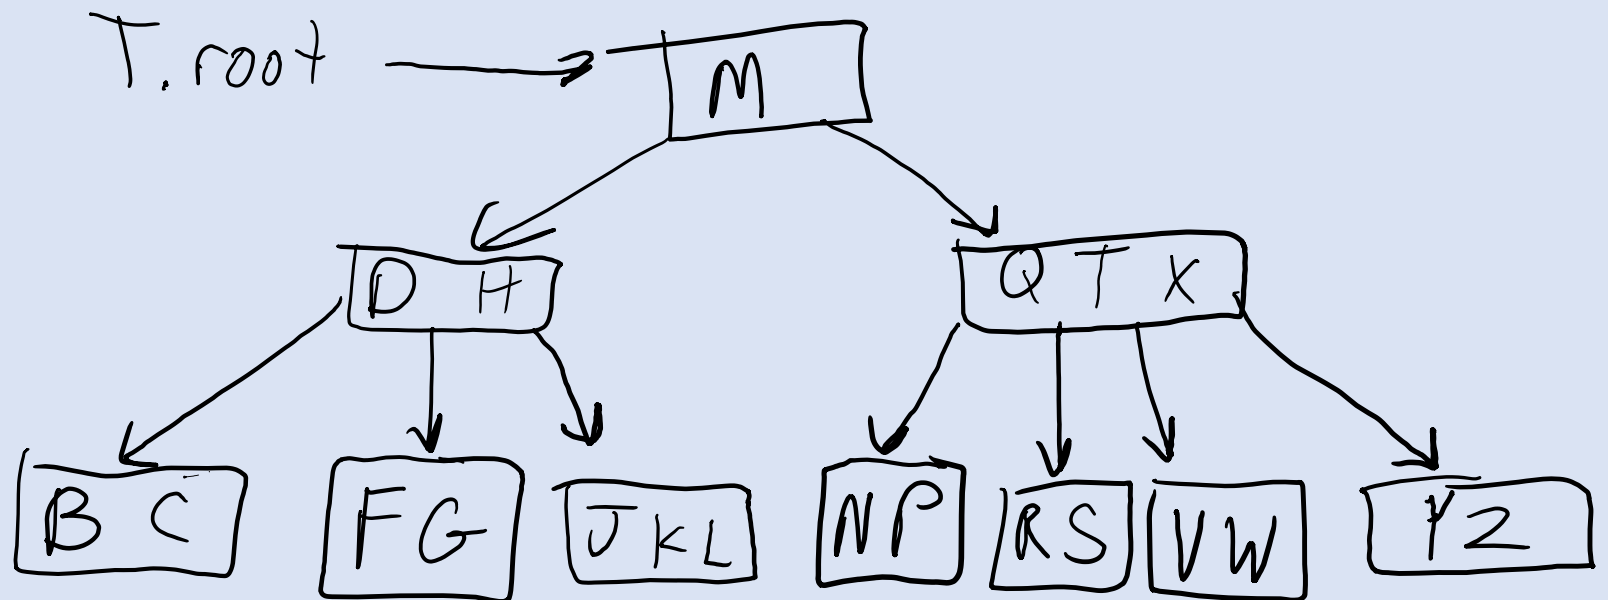


B-Trees

- Balanced Search Trees
- One con with a Binary Search Tree (BST) is that we can potentially have our algorithms run in linear time rather than height level time.
- B-Trees nodes have many children (few to even thousands!)
- This data structure is primarily used in disks and other direct-access secondary storage devices.
- Database systems use B-Trees or even variants.
- The height of a B-tree is $O(\lg n)$

B-Tree Sample

- In this sample, the keys are letters from the alphabet
- Each node has $x.n$ keys and $x.n + 1$ children
- How does a search work for the letter P?



Data Structures on Secondary Storage

- Primary memory (main memory) consists of silicon memory chips.
- Secondary storage consists of magnetic storage
 - Tapes
 - Disks
- Disks are cheaper and have higher capacity than the main memory.
- Disks are slower than main memory due to motion mechanical components.

Disk Drive

- The average access time for the disk ranges from 8 to 11 milliseconds.
- The average access time in main memory is about 50 nanoseconds!
- Information on a disk is divided into pages which range from 2^{11} – 2^{14} bytes.
- Each disk reads and/or writes on a single or multiple pages.

B-Tree Applications

- The Whole B-Trees do not fit in the main memory!!!
- Operating Systems copies the pages from the disks into main memory. After performing tasks, the operating system writes back to the respective pages that were modified.

x = a pointer to some object

DISK-READ(x)


operations that access and/or modify attributes of x

DISK-WRITE(x)

other operations that access but do not modify attributes of x

B-Tree Example

- Branch Factor = 1001
- Height = 2



- B-Trees often have branching factors ranging from 50 -2000
- Root node is permanently in main memory in order to find any key with at most two disk accesses.

The Official B-Tree Definition


- A B-Tree is a rooted tree (where T.root is the root) with the following properties:
 1. Every node x has the following attributes
 - a) $x.n$ is the number of keys currently stored in x
 - b) The keys $x.key_1, x.key_2, \dots, x.key_{x.n}$ such that
$$x.key_1 \leq x.key_2 \leq \dots \leq x.key_{x.n}$$
 - c) $x.leaf$ – a Boolean value which is true if x is a leaf and false if x is an internal node
 2. Each internal node x has $x.n + 1$ pointers $x.c_1, x.c_2, \dots, x.c_{x.n+1}$ to its children. If x is a leaf then the pointers are undefined.
 3. If k_i is any key stored in the subtree with root $x.c_i$ then:
$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

The official B-Tree Definition Continued

- All leaves have the same depth, which is the tree high h .
- The B-Tree has a minimum degree t (where t is an interger $t \geq 2$):
 - Every node other than the root must have $\geq t - 1$ keys and $\geq t$ children; if B-tree is nonempty, then the root has at least one key
 - Every node has $\leq 2t - 1$ keys and $\leq 2t$ children
A node is considered full if it has $2t - 1$ keys inserted.

Interesting Theorem About Height in B-Trees

- Theorem: if $n \geq 1$, then for any n -key B-tree T of height h and minimum degree t ,



Theorem cont.

- $n \geq 1 + (t - 1) \sum_{i=1}^h 2t^{i-1}$

$$= 1 + 2(t - 1) \sum_{i=1}^h t^{i-1}$$

$$= 1 + 2(t - 1) \frac{t^h - 1}{t - 1} = 2t^h - 1$$

$$t^h \leq \frac{n + 1}{2}$$

$$h \leq \log_t \frac{n + 1}{2}$$

$$h = O(\log n)$$

Operations We Will Observe for B-Trees

- B-Tree-Search
- B-Tree-Create
- B-Tree-Insert
- B-Tree-Delete

B-Tree Search

B-Tree-Search(x, k)

$i = 1$

while $i \leq x.n$ and $k > x$

$i = i + 1$

if $i \leq x.n$ and $k == x.key_i$

return (x, i)

else if $x.leaf == \text{True}$

return NULL

else DISK-READ($x.c_i$)

return B-Tree-Search($x.c_i, k$)

RT: $O(t \log_t n)$

B-Tree-Create

- Creating an empty tree with root node

B-Tree-Create(T)

x = Allocate-Node()

x.leaf = True

x.n = 0

Disk-Write(x)

T.root = x


B-Tree Insert Operations

- The insert operation has 3 functions/methods we need to understand
- B-Tree-Split-Child(x,i)
- B-Tree-Insert(T,k)
- B-Tree-Insert-Nonfull(x,k)

Insert Operation and Overall Goal

- Search for a leaf where to put new key
- Inserting into an existing leaf node
 - Cannot create a new leaf
- If the leaf node is full, then split around the median key
- The overall goal is to insert the key while maintaining B-Tree rules. As the algorithm traverses down the tree, it splits each full node along the way, including the leaf.

Splitting a Node

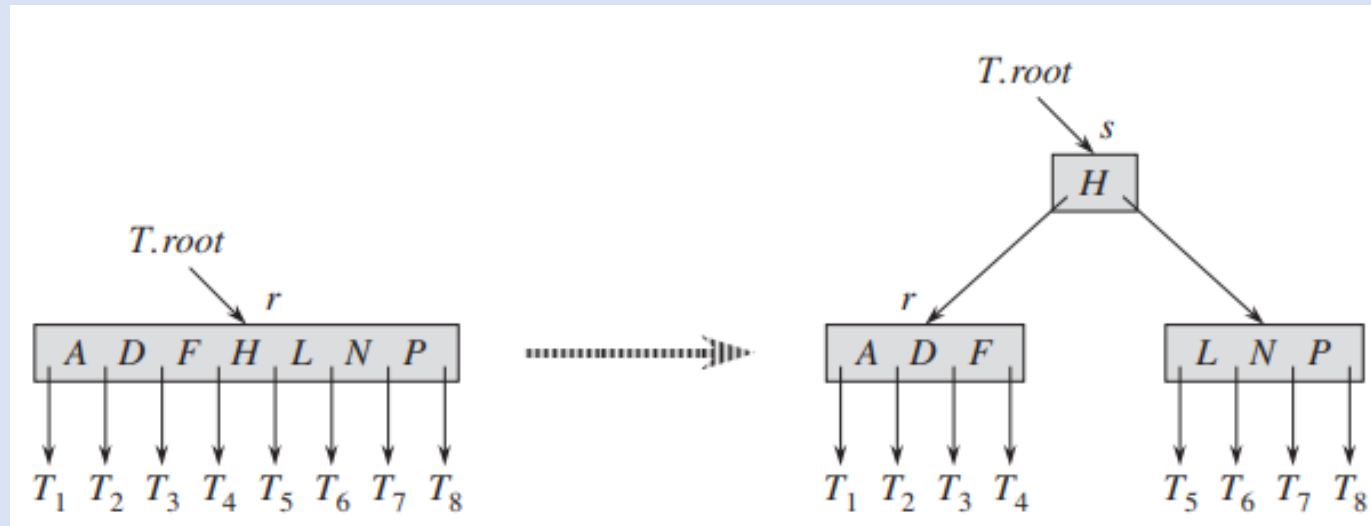


B-TREE-SPLIT-CHILD(x, i)

```
1   $z = \text{ALLOCATE-NODE}()$ 
2   $y = x.c_i$ 
3   $z.\text{leaf} = y.\text{leaf}$ 
4   $z.n = t - 1$ 
5  for  $j = 1$  to  $t - 1$ 
6       $z.\text{key}_j = y.\text{key}_{j+t}$ 
7  if not  $y.\text{leaf}$ 
8      for  $j = 1$  to  $t$ 
9           $z.c_j = y.c_{j+t}$ 
10  $y.n = t - 1$ 
11 for  $j = x.n + 1$  downto  $i + 1$ 
12      $x.c_{j+1} = x.c_j$ 
13  $x.c_{i+1} = z$ 
14 for  $j = x.n$  downto  $i$ 
15      $x.\text{key}_{j+1} = x.\text{key}_j$ 
16  $x.\text{key}_i = y.\text{key}_t$ 
17  $x.n = x.n + 1$ 
18  $\text{DISK-WRITE}(y)$ 
19  $\text{DISK-WRITE}(z)$ 
20  $\text{DISK-WRITE}(x)$ 
```

B-Tree-Insert()

- $t = 4$
range of keys 3-7



- If root node is full, then split the root and new node will become the root

B-TREE-INSERT(T, k)

```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{ALLOCATE-NODE}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
```

B-TREE-INSERT-NONFULL(x, k)

```
1   $i = x.n$ 
2  if  $x.leaf$ 
3      while  $i \geq 1$  and  $k < x.key_i$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$ 
7       $x.n = x.n + 1$ 
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$ 
10       $i = i - 1$ 
11       $i = i + 1$ 
12      DISK-READ( $x.c_i$ )
13      if  $x.c_i.n == 2t - 1$ 
14          B-TREE-SPLIT-CHILD( $x, i$ )
15          if  $k > x.key_i$ 
16               $i = i + 1$ 
17          B-TREE-INSERT-NONFULL( $x.c_i, k$ )
```

Insertion Examples

RT for Insert

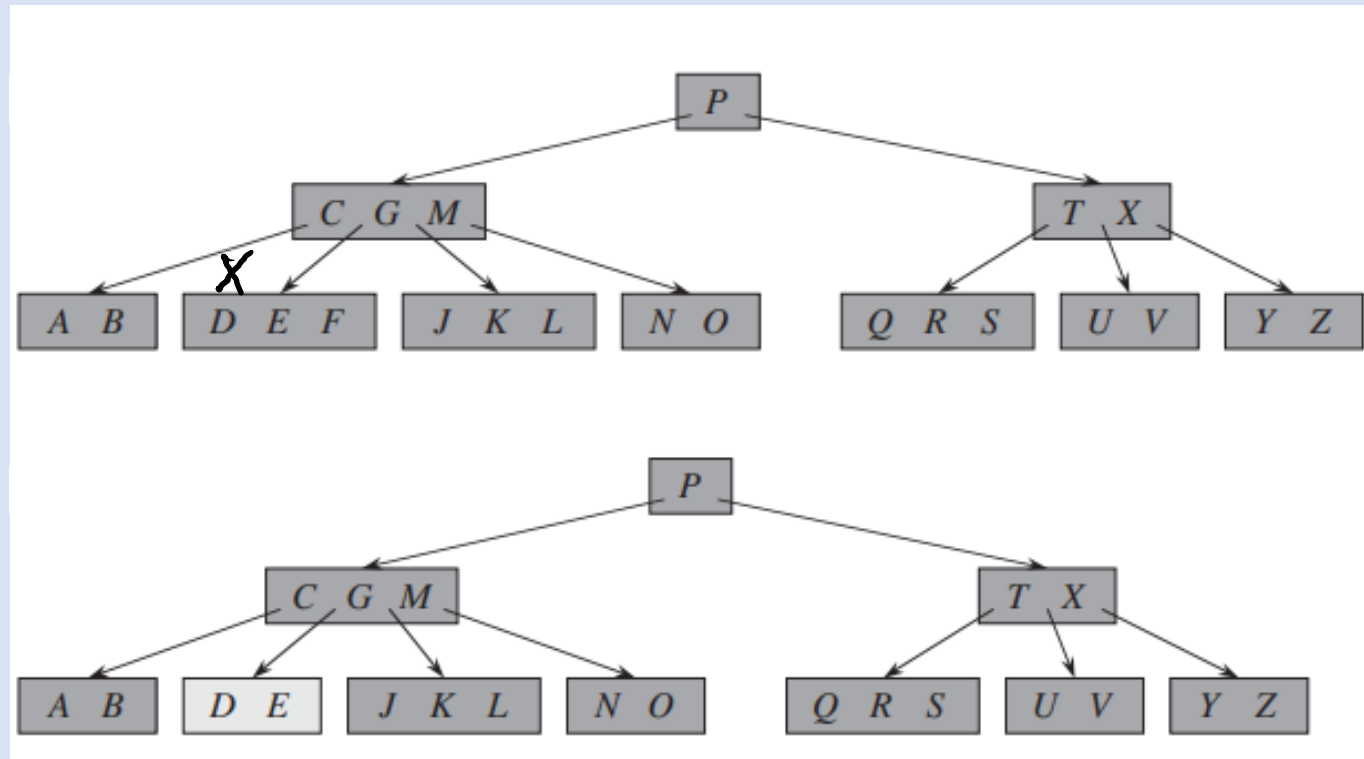
- $O(th) = O(t \log_t n)$

B-Tree-Delete()

- Important things to remember!
 - When a key is removed, we must rearrange the node's children!
 - Any node (EXCEPT FOR THE ROOT) cannot have fewer than $t - 1$ keys
 - The algorithm deletes a key k from the subtree rooted at x
 - Something to consider: When delete is called on a node, we should guarantee that the number of keys in x is greater than or equal to t .
- The overall objective is to remove a key while maintaining the B-tree properties.
- There are 3 rules to consider when deleting from the B-Tree

Rule 1

- If the key k is part of a leaf node x , then just delete the key.
 $t = 3$




Delete F

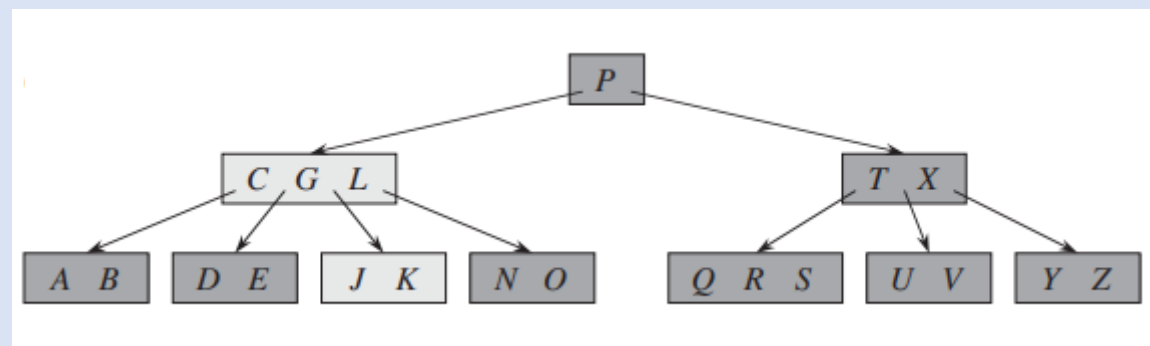
Rule 2a

- If the key k belongs to an internal node x .
- If the child y that precedes k in a node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' and replace k by k' in x .

$t = 3$



Delete M




Rule 2b

- If the key k belongs to an internal node.
- If y has fewer than t keys, then, symmetrically, examine the child z that follows k in node x . If z has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k' and replace k by k' in x .


Rule 2c

- If the key k belongs to an internal node x .
- Otherwise, if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then free z and recursively delete k from y .

$t = 3$



Delete G





Rule 3a

- If the key k is not part of the internal node x , take $x.c_i$ the root of the subtree that must contain k (if k is in the tree). If $x.c_i$ has only $t-1$ keys, then use 3a or 3b to guarantee we descend to a node with greater than or equal to t keys.

If $x.c_i$ has an immediate sibling with greater than or equal to t keys, then give $x.c_i$ an extra key by:

- Moving a key from x to $x.c_i$
- Moving a key from $x.c_i$'s immediate left or right sibling up x
- Moving the appropriate child pointer from the sibling into $x.c_i$


$t = 3$
Delete B



Rule 3b

- If both $x.c_i$'s immediate sibling have $t - 1$ keys, merge $x.c_i$ with one sibling, which involves moving a key from x down into the new merged node to become the median for that node

Delete D



RT for Delete

- RT is $O(th) = O(t \log_t n)$

Delete Examples